

IMDB Movie Sentiment Analysis

In this notebook, we will use the IMDB movie reviews dataset, provided on Kaggle (<https://www.kaggle.com/c/word2vec-nlp-tutorial/data> (<https://www.kaggle.com/c/word2vec-nlp-tutorial/data>)), to build a sentiment analysis model using NLP.

Rather than using the Word2Vec model (which provided little to no gains in our score), this time we will be using the Doc2Vec model which will consider the order of words in our review as well, giving us not only each word's average feature vector, but also the whole review's average feature vector as well. Through this, we hope to improve our score.

Importing the Data

```
In [404]: import numpy as np
import pandas as pd # to read the csv datasets and import them into python
```

```
In [405]: data_path = "data" # path to the data folder in your directory
# import the labeled and unlabeled training data to train our model
train_1 = pd.read_csv(data_path + "/labeledTrainData.tsv", header=0, delimiter
="\t", quoting=3)
train_2 = pd.read_csv(data_path + "/imdb_master.csv", encoding="latin-1")
```

```
In [406]: train_2 = train_2.drop(["Unnamed: 0", "type", "file"], axis=1)
train_2.columns = ["review", "sentiment"]
```

```
In [407]: train_2 = train_2[train_2.sentiment != 'unsup']
train_2['sentiment'] = train_2['sentiment'].map({'pos': 1, 'neg': 0})
```

```
In [408]: train = pd.concat([train_1, train_2]).reset_index(drop=True)
```

Preprocessing

```
In [409]: from bs4 import BeautifulSoup # to get rid of the HTML tags in the reviews
import re # to remove punctuations and numericals from the review

from nltk.corpus import stopwords # to remove the stop words in our reviews and obtain our tokenizer
from nltk.stem import WordNetLemmatizer # to lemmatize our reviews
from nltk import word_tokenize
import nltk.data
nltk.download()
lemmatizer = WordNetLemmatizer()
stop_words = set(stopwords.words("english"))
```

showing info https://raw.githubusercontent.com/nltk/nltk_data/gh-pages/index.xml

```
In [533]: def preprocess_review(unclean_review):
        """
        Function that takes a single unclean review from the original dataset
        and returns a cleaned and preprocessed version of it.
        Input: string: an uncleaned review from the dataset
        Output: string: cleaned and preprocessed review
        """
        # removes the HTML tags in the review
        untagged_review = BeautifulSoup(unclean_review).get_text()
        # removes everything not in A-Z or a-z and replaces it with a space
        letter_only_review = re.sub("[^a-zA-Z]", " ", untagged_review)
        # converting everything to lowercase
        letter_only_review = letter_only_review.lower()
        # converting everything to tokenized words
        tokenized_review = word_tokenize(letter_only_review)
        # taking only the lemmatized words
        tokens = list(map(lemmatizer.lemmatize, tokenized_review))
        # only accepting the verbs from those lemmatized words to normalize process
        lemmatized_tokens = list(map(lambda x: lemmatizer.lemmatize(x, "v"), tokens))
        # removing all the stop words in the review
        result = [t for t in lemmatized_tokens if not t in stop_words]
        return result
```

Training the Model

```
In [411]: # clean up the reviews
x_train = []
for i in range(len(train["review"])):
    if i % 10000 == 0:
        print(i)
        x_train.append(preprocess_review(train["review"][i]))

y_train = train.sentiment.values
```

```
0
10000
20000
30000
40000
50000
60000
70000
```

```
In [412]: from sklearn.model_selection import train_test_split
from keras.preprocessing.text import Tokenizer
from keras.preprocessing.sequence import pad_sequences
from keras.layers import Dense , Input , LSTM , Embedding, Dropout , Activation, GRU, Flatten
from keras.layers import Bidirectional, GlobalMaxPool1D
from keras.models import Model, Sequential
from keras.layers import Convolution1D
from keras import initializers, regularizers, constraints, optimizers, layers

max_features = 5000 # we will keep the 5000 most freq words for our model
tokenizer = Tokenizer(num_words=max_features)
tokenizer.fit_on_texts(x_train)
train_tokens = tokenizer.texts_to_sequences(x_train)

max_length = 130
x_train = pad_sequences(train_tokens, maxlen=max_length)

x_train, x_test, y_train, y_test = train_test_split(x_train, y_train, test_size = 0.1, shuffle=True)
embedding_size = 128
model = Sequential()
model.add(Embedding(max_features, embedding_size))
model.add(Bidirectional(LSTM(32, return_sequences = True)))
model.add(GlobalMaxPool1D())
model.add(Dense(20, activation="relu"))
model.add(Dropout(0.05))
model.add(Dense(1, activation="sigmoid"))
model.compile(loss='binary_crossentropy', optimizer='adam', metrics=['accuracy'])

batch_size = 100
epochs = 8
history = model.fit(x=x_train, y=y_train, validation_data=(x_test, y_test),
                    batch_size=batch_size, epochs=epochs)
```

C:\Users\yasoo\Anaconda3\envs\walmart1\lib\site-packages\tensorflow_core\python\framework\indexed_slices.py:433: UserWarning: Converting sparse IndexedSlices to a dense Tensor of unknown shape. This may consume a large amount of memory.

"Converting sparse IndexedSlices to a dense Tensor of unknown shape. "

Train on 67500 samples, validate on 7500 samples

Epoch 1/8

67500/67500 [=====] - 111s 2ms/step - loss: 0.3199 - accuracy: 0.8636 - val_loss: 0.2497 - val_accuracy: 0.9003

Epoch 2/8

67500/67500 [=====] - 113s 2ms/step - loss: 0.2103 - accuracy: 0.9207 - val_loss: 0.2285 - val_accuracy: 0.9123

Epoch 3/8

67500/67500 [=====] - 117s 2ms/step - loss: 0.1586 - accuracy: 0.9443 - val_loss: 0.2161 - val_accuracy: 0.9228

Epoch 4/8

67500/67500 [=====] - 116s 2ms/step - loss: 0.1177 - accuracy: 0.9603 - val_loss: 0.2189 - val_accuracy: 0.9296

Epoch 5/8

67500/67500 [=====] - 116s 2ms/step - loss: 0.0820 - accuracy: 0.9739 - val_loss: 0.2325 - val_accuracy: 0.9335

Epoch 6/8

67500/67500 [=====] - 119s 2ms/step - loss: 0.0584 - accuracy: 0.9828 - val_loss: 0.2360 - val_accuracy: 0.9349

Epoch 7/8

67500/67500 [=====] - 118s 2ms/step - loss: 0.0422 - accuracy: 0.9876 - val_loss: 0.2843 - val_accuracy: 0.9364

Epoch 8/8

67500/67500 [=====] - 119s 2ms/step - loss: 0.0332 - accuracy: 0.9899 - val_loss: 0.2969 - val_accuracy: 0.9348

Output Model Results

```
In [413]: # import the test data to evaluate our model
result_data = pd.read_csv(data_path + "/testData.tsv", header=0, delimiter="\t", quoting=3)
```

```
In [414]: result_data["review"] = result_data.review.apply(lambda x: preprocess_review(x))
```

```
In [415]: result_data["sentiment"] = result_data["id"].map(lambda x: 1 if int(x.strip(' ').split("_")[1]) >= 5 else 0)
```

```
In [416]: y_result = result_data["sentiment"]
reviews = result_data["review"]
tokenized = tokenizer.texts_to_sequences(reviews)
x_result = pad_sequences(tokenized, maxlen=max_length)
```

```
In [417]: predictions = model.predict(x_result)
```

```
In [418]: y_result_pred = (predictions > 0.5)
```

```
In [419]: # to obtain the area under ROC score
from sklearn.metrics import f1_score, confusion_matrix
print('F1-score: {0}'.format(f1_score(y_result_pred, y_result)))
print('Confusion matrix:')
confusion_matrix(y_result_pred, y_result)
```

F1-score: 0.9769425777634704

Confusion matrix:

```
Out[419]: array([[12288,   361],
                 [  212, 12139]], dtype=int64)
```

```
In [420]: y_result_pred = y_result_pred.astype(int)
```

```
In [421]: y_result_pred = [item for items in y_result_pred for item in items]
```

```
In [422]: output = pd.DataFrame(data={"id":test_data["id"], "sentiment":y_result_pred} )
output.to_csv( "cnn_model.csv", index=False, quoting=3 )
```

```
In [423]: output
```

```
Out[423]:
```

	id	sentiment
0	"12311_10"	1
1	"8348_2"	0
2	"5828_4"	0
3	"7186_2"	0
4	"12128_7"	1
...
24995	"2155_10"	1
24996	"59_10"	1
24997	"2531_1"	0
24998	"7772_8"	1
24999	"11465_10"	1

25000 rows × 2 columns

```
In [ ]:
```

Visualizing the Model

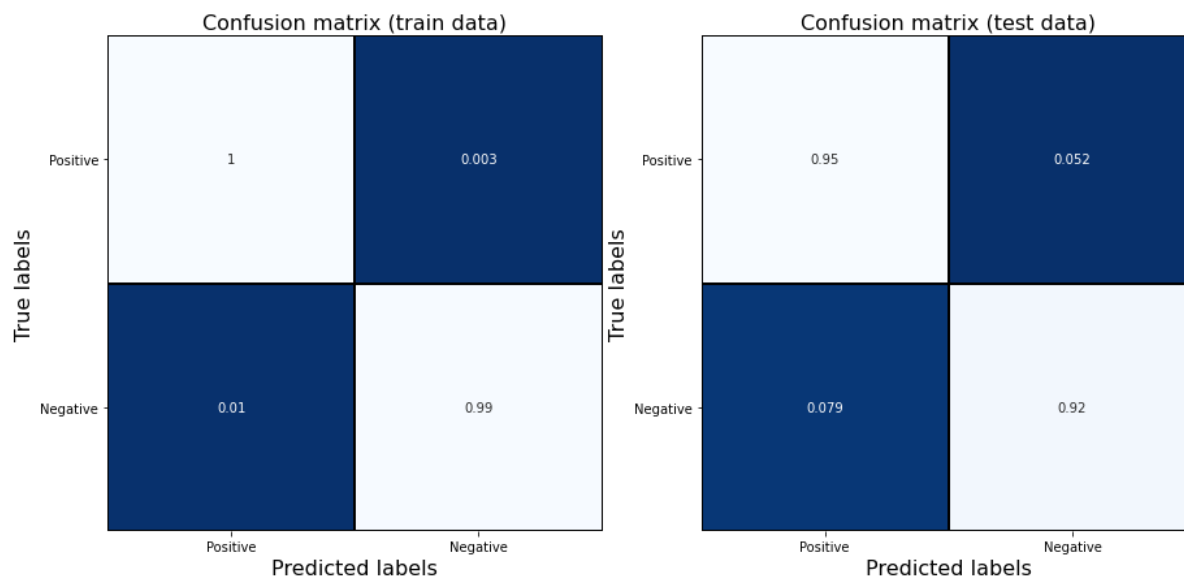
In this part, we will visualize the confusion matrix and the model accuracy/loss curves, using guidance from [here](https://www.kaggle.com/alexcherniuk/imdb-review-word2vec-bilstm-99-acc) (<https://www.kaggle.com/alexcherniuk/imdb-review-word2vec-bilstm-99-acc>)

```
In [424]: y_train_predictions = model.predict_classes(x_train)
y_test_predictions = model.predict_classes(x_test)
```

```
In [425]: # library to plot the confusion matrix and our curves
import matplotlib.pyplot as plt
from sklearn.metrics import accuracy_score, confusion_matrix
import seaborn as sb

# function to plot the confusion matrix
def plot_confusion_matrix(y_true, y_pred, ax, class_names, vmax=None,
                          normed=True, title='Confusion matrix'):
    matrix = confusion_matrix(y_true, y_pred)
    if normed:
        matrix = matrix.astype('float') / matrix.sum(axis=1)[:, np.newaxis]
    sb.heatmap(matrix, annot=True, square=True, ax=ax,
               cmap=plt.cm.Blues_r, cbar=False, linecolor='black',
               linewidths=1, xticklabels=class_names)
    ax.set_title(title, y=1, fontsize=16)
    ax.set_ylabel('True labels', fontsize=16)
    ax.set_xlabel('Predicted labels', y=1, fontsize=16)
    ax.set_yticklabels(class_names, rotation=0)
```

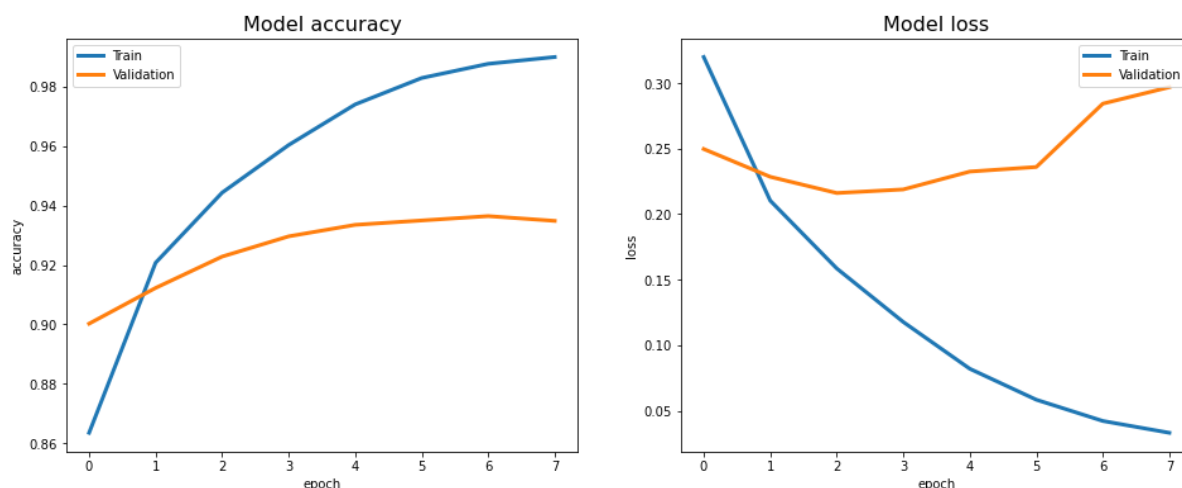
```
In [426]: fig, (axis1, axis2) = plt.subplots(nrows=1, ncols=2, figsize=(15, 15))
plot_confusion_matrix(y_train, y_train_predictions, ax=axis1,
                      title='Confusion matrix (train data)',
                      class_names=['Positive', 'Negative'])
plot_confusion_matrix(y_test, y_test_predictions, ax=axis2,
                      title='Confusion matrix (test data)',
                      class_names=['Positive', 'Negative'])
```



```
In [427]: fig, (axis1, axis2) = plt.subplots(nrows=1, ncols=2, figsize=(16,6))

# summarize history for accuracy
axis1.plot(history.history['accuracy'], label='Train', linewidth=3)
axis1.plot(history.history['val_accuracy'], label='Validation', linewidth=3)
axis1.set_title('Model accuracy', fontsize=16)
axis1.set_ylabel('accuracy')
axis1.set_xlabel('epoch')
axis1.legend(loc='upper left')

# summarize history for loss
axis2.plot(history.history['loss'], label='Train', linewidth=3)
axis2.plot(history.history['val_loss'], label='Validation', linewidth=3)
axis2.set_title('Model loss', fontsize=16)
axis2.set_ylabel('loss')
axis2.set_xlabel('epoch')
axis2.legend(loc='upper right')
plt.show()
```



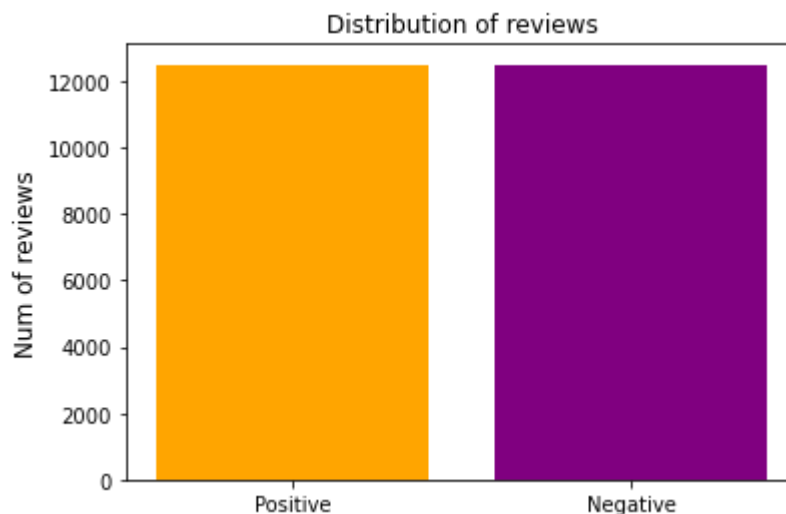
In []:

Analysis

1) How many movie reviews are positive and how many are negative in labeledTrainData.tsv? Do we have balance between the classes?


```
In [453]: positive = train_1[train_1.sentiment == 1]
negative = train_1[train_1.sentiment == 0]
x = ["Positive", "Negative"]
y = [len(positive), len(negative)]
plt.bar(x, y, color=["orange", "purple"])
plt.title("Distribution of reviews")
plt.ylabel("Num of reviews", fontsize=12)
```

```
Out[453]: Text(0, 0.5, 'Num of reviews')
```



As we can observe, both the positive and negative reviews are equally balanced (12,500 each) in the labeledTrainData.tsv

2) What is the average length of all the cleaned up reviews (string length)?

```
In [562]: def clean_frame(dataframe):
cleaned_reviews = []
i = 0
for review in dataframe["review"]:
    if i % 10000 == 0:
        print(i)
    cleaned_reviews.append(preprocess_review(review))
    i += 1
return cleaned_reviews
```

```
In [ ]: cleaned_reviews = clean_frame(train_1)
```

```
In [536]: import scipy, statistics
review_lengths = []
for review in cleaned_reviews:
    review_lengths.append(len(review))
average = np.mean(review_lengths)
mode = statistics.mode(review_lengths)
median = np.median(review_lengths)
```

```
In [547]: fig, ax = plt.subplots(figsize=(15, 4))
sb.distplot(review_lengths, bins=2000, ax=ax, color="orange")
ax.set_xlim(left=0, right=np.percentile(review_lengths, 95))
ax.set_xlabel("Words in a review", fontsize=16)
ax.set_ylabel("Distribution of words", fontsize=16)
plt.axvline(x=average, ymin=0, ymax=1, label="Average", color="red")
plt.axvline(x=mode, ymin=0, ymax=1, label="Mode", color="blue")
plt.axvline(x=median, ymin=0, ymax=1, label="Median", color="green")
plt.legend()
print("Average: {}, Mode: {}, Median: {}".format(average, mode, median))
```

Average: 122.18704, Mode: 66, Median: 91.0



What does the word cloud for positive review words look like? Negative review words?

```
In [563]: positive["review"]
```

```
Out[563]: 0      "With all this stuff going down at the moment ..."
1      "\"The Classic War of the Worlds\" by Timothy ..."
4      "Superbly trashy and wondrously unpretentious ..."
5      "I dont know why people think this is such a b..."
9      "<br /><br />This movie is full of references...."
...
24987  "First off, I'd like to make a correction on a..."
24988  "While originally reluctant to jump on the ban..."
24989  "I heard about this movie when watching VH1's ..."
24990  "I've never been huge on IMAX films. They're c..."
24999  "I saw this movie as a child and it broke my h..."
Name: review, Length: 12500, dtype: object
```

```
In [584]: clean_positive = clean_frame(positive)
clean_negative = clean_frame(negative)
```

```
0
10000
0
10000
```

```
In [586]: pos_word_list = ""
          for review in clean_positive:
              pos_word_list += " ".join(review)

          neg_word_list = ""
          for review in clean_negative:
              neg_word_list += " ".join(review)
```

```
In [589]: from wordcloud import WordCloud

def generate_wordcloud(words, heading):
    """
    Function that generates a word cloud for us using the words we supply to it
    """
    # instantiate the word cloud
    wordcloud = WordCloud(stopwords=stop_words, max_words=100, max_font_size=40, scale=3).generate(words)
    # generate the figure
    fig = plt.figure(figsize=(15, 15))
    plt.title(heading, fontsize=15)
    plt.imshow(wordcloud)
    plt.show()

generate_wordcloud(pos_word_list, "Positive")
generate_wordcloud(neg_word_list, "Negative")
```

